



ASSIGNMENT 1

Tecnologias de Bases de Dados

April 2023

GROUP D

Mariana Monteiro

up202003480@fe.up.pt

Vasco Gomes

up201906617@fe.up.pt

Índice

Constraints Y	3
Constraints Z	4
Query 1	5
1) SQL query.....	5
2) Answer.....	5
3) Three execution plans.....	5
4) Execution time.....	6
Query 2	7
1) SQL query.....	7
2) Answer.....	7
3) Three execution plans.....	7
4) Execution time.....	8
Query 3	9
a	9
1) SQL query	9
2) Answer.....	9
3) Three execution plans	10
4) Execution time.....	10
b	11
1) SQL query	11
2) Answer.....	11
3) Three execution plans	12
4) Execution time.....	12
Query 4	13
1) SQL query.....	13
2) Answer.....	13
3) Three execution plans.....	13
4) Execution time	15
Query 5	16
1) SQL query.....	16
a.	16
b.	16
2) Answer.....	16
3) Three execution plans.....	17
a.	17
b.	17
4) Execution time.....	17
Query 6	18

1) SQL query.....	18
2) Answer.....	18
3) Three execution plans.....	18
4) Execution time	19

Constraints Y

Defining primary key and foreign key constraints on MySQL tables can significantly improve query performance.

When a primary key is defined on a table, an index is automatically created on that column. This index enables the database engine to easily locate specific rows in the table, which can greatly reduce query execution time.

Similarly, when a foreign key is defined, it establishes a relationship between two tables that can be used to join them together efficiently.

By utilizing these constraints, the database engine can avoid expensive table scans and instead use indexes to quickly retrieve the necessary data and enhance some queries performance.

```
ALTER TABLE ydocentes ADD CONSTRAINT ydocentes_pk PRIMARY KEY (nr);  
ALTER TABLE zdocentes ADD CONSTRAINT zdocentes_pk PRIMARY KEY (nr);
```

```
ALTER TABLE yucs ADD CONSTRAINT yucs_pk PRIMARY KEY (codigo);  
ALTER TABLE zucs ADD CONSTRAINT zucs_pk PRIMARY KEY (codigo);
```

```
ALTER TABLE yocorrencias ADD CONSTRAINT yocorrencias_pk PRIMARY KEY (codigo,  
ano_letivo, periodo);  
ALTER TABLE yocorrencias ADD CONSTRAINT yocorrencias_fk FOREIGN KEY(codigo)  
REFERENCES yucs(codigo);  
ALTER TABLE zocorrencias ADD CONSTRAINT zocorrencias_pk PRIMARY KEY (codigo,  
ano_letivo, periodo);  
ALTER TABLE zocorrencias ADD CONSTRAINT zocorrencias_fk FOREIGN KEY(codigo)  
REFERENCES zucs(codigo);
```

```
ALTER TABLE ytiposaula ADD CONSTRAINT ytiposaula_pk PRIMARY KEY(id);  
ALTER TABLE ytiposaula ADD CONSTRAINT ytiposaula_fk FOREIGN KEY (ano_letivo,  
periodo, codigo) REFERENCES yocorrencias(ano_letivo, periodo, codigo);  
ALTER TABLE ztiposaula ADD CONSTRAINT ztiposaula_pk PRIMARY KEY(id);  
ALTER TABLE ztiposaula ADD CONSTRAINT ztiposaula_fk FOREIGN KEY (ano_letivo,  
periodo, codigo) REFERENCES zocorrencias(ano_letivo, periodo, codigo);
```

```
ALTER TABLE ydsd ADD CONSTRAINT ydsd_pk PRIMARY KEY (nr,id);  
ALTER TABLE ydsd ADD CONSTRAINT ydsd_fk1 FOREIGN KEY (nr) REFERENCES  
ydocentes(nr);  
ALTER TABLE ydsd ADD CONSTRAINT ydsd_fk2 FOREIGN KEY (id) REFERENCES  
ytiposaula(id);  
ALTER TABLE zdsd ADD CONSTRAINT zdsd_pk PRIMARY KEY (nr,id);  
ALTER TABLE zdsd ADD CONSTRAINT zdsd_fk1 FOREIGN KEY (nr) REFERENCES  
zdocentes(nr);  
ALTER TABLE zdsd ADD CONSTRAINT zdsd_fk2 FOREIGN KEY (id) REFERENCES  
ztiposaula(id);
```

Constraints Z

Overall, we consider indexing columns that are frequently used in the WHERE clause of our queries. By indexing these columns, we can improve the performance of queries that filter rows based on those columns.

We chose to index **curso** and **designacao** on the table **zucs**, **curso** on the table **zucs**, **ano_letivo** on the table **zocorrencias** and **ano_letivo** on the table **ztiposaula**.

```
CREATE INDEX ucs_idx ON zucs (curso, designacao);  
CREATE INDEX ucs_idx_curso ON zucs (curso);  
CREATE INDEX ocorrencias_idx ON zocorrencias (ano_letivo);  
CREATE INDEX tiposaula_idx_anoletivo ON ztiposaula (ano_letivo);
```

Query 1 - Show the codigo, designacao, ano_letivo, inscritos, tipo, and turnos for the course 'Bases de Dados' of the program 275.

1) the SQL query;

```
SELECT xucs.codigo, xucs.designacao, xocorrencias.ano_letivo,
xocorrencias.inscritos,xtiposaula.tipo,xtiposaula.turnos
FROM xucs
JOIN xocorrencias ON xucs.codigo=xocorrencias.codigo
JOIN xtiposaula ON xocorrencias.ano_letivo=xtiposaula.ano_letivo
AND xocorrencias.periodo= xtiposaula.periodo
AND xocorrencias.codigo=xtiposaula.codigo
WHERE curso='275' AND designacao='Bases de Dados';
```

2) the answer;

	CODIGO	DESIGNACAO	ANO_LETIVO	INSCRITOS	TIPO	TURNOS
1	EIC3106	Bases de Dados	2003/2004	92	T	1
2	EIC3106	Bases de Dados	2003/2004	92	TP	4
3	EIC3106	Bases de Dados	2004/2005	114	T	1
4	EIC3106	Bases de Dados	2004/2005	114	TP	4
5	EIC3111	Bases de Dados	2005/2006	(null)	T	1
6	EIC3111	Bases de Dados	2005/2006	(null)	TP	6

3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	642
HASH JOIN			1	642
Access Predicates				
AND				
XOCORRENCIAS.ANO_LETIVO=XTIPOSADULA.ANO_LETIVO				
XOCORRENCIAS.PERIODO=XTIPOSADULA.PERIODO				
XOCORRENCIAS.CODIGO=XTIPOSADULA.CODIGO				
XUCS.CODIGO=XOCORRENCIAS.CODIGO				
MERGE JOIN		CARTESIAN	4768	49
TABLE ACCESS	XUCS	FULL	1	13
Filter Predicates				
AND				
XUCS.DESIGNACAO='Bases de Dados'				
XUCS.CURSO=275				
BUFFER		SORT	21019	36
TABLE ACCESS	XTIPOSADULA	FULL	21019	36
TABLE ACCESS	XOCORRENCIAS	FULL	21747	593
Other XML				

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	55
HASH JOIN			1	55
Access Predicates				
AND				
YOCORRENCIAS.ANO_LETIVO=YTIPOSADULA.ANO_LETIVO				
YOCORRENCIAS.PERIODO=YTIPOSADULA.PERIODO				
YOCORRENCIAS.CODIGO=YTIPOSADULA.CODIGO				
NESTED LOOPS			1	19
NESTED LOOPS			5	19
TABLE ACCESS	YUCS	FULL	1	13
Filter Predicates				
AND				
YUCS.DESIGNACAO='Bases de Dados'				
YUCS.CURSO=275				
INDEX	YOCORRENCIAS_PK	RANGE SCAN	5	1
Access Predicates				
YUCS.CODIGO=YOCORRENCIAS.CODIGO				
TABLE ACCESS	YOCORRENCIAS	BY INDEX ROWID	5	6
TABLE ACCESS	YTIPOSADULA	FULL	21019	36
Other XML				

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	44
HASH JOIN			1	44
Access Predicates				
AND				
ZOCORRENCIAS.ANO_LETIVO=ZTIPOSAULA.ANO_LETIVO				
ZOCORRENCIAS.PERIODO=ZTIPOSAULA.PERIODO				
ZOCORRENCIAS.CODIGO=ZTIPOSAULA.CODIGO				
NESTED LOOPS			1	8
NESTED LOOPS			5	8
TABLE ACCESS	ZUCS	BY INDEX ROWID BATCHED	1	2
INDEX	UCS_IDX	RANGE SCAN	1	1
Access Predicates				
AND				
ZUCS.CURSO=275				
ZUCS.DESIGNACAO='Bases de Dados'				
INDEX	ZOCORRENCIAS_PK	RANGE SCAN	5	1
Access Predicates				
ZUCS.CODIGO=ZOCORRENCIAS.CODIGO				
TABLE ACCESS	ZOCORRENCIAS	BY INDEX ROWID	5	6
TABLE ACCESS	ZTIPOSAULA	FULL	21019	36
Other XML				

The usage of the foreign keys codigo and ano_letivo improved the cost of the JOIN between yocorrencias and yucs and the JOIN between yocorrencias and ytiposaula.

The usage of the index ucs_index improved the cost of filtering the curso and designacao in environment z.

4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment X	78
Environment Y	37
Environment Z	34

Query 2 - How many class hours of each type did the program 233 planned in year 2004/2005?

1) the SQL query;

```
SELECT tipo, SUM(horas) as hours_planned
FROM xtiposaula s JOIN xdsd d ON d.id=s.id
JOIN xocorrencias o ON o.codigo=s.codigo AND o.ano_letivo=s.ano_letivo AND
o.periodo=s.periodo
JOIN xucs u ON o.codigo=u.codigo
WHERE u.curso=233 AND o.ano_letivo='2004/2005'
GROUP BY s.tipo;
```

2) the answer;

	TIPO	HOURS_PLANNED
1	P	571,5
2	TP	697,5
3	T	298

3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort;

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			5	670
HASH		GROUP BY	5	670
HASH JOIN			894	669
Access Predicates				
D.ID=S.ID				
HASH JOIN			552	642
Access Predicates				
AND				
O.CODIGO=S.CODIGO				
O.ANO_LETIVO=S.ANO_LETIVO				
O.PERIODO=S.PERIODO				
HASH JOIN			552	606
Access Predicates				
O.CODIGO=U.CODIGO				
TABLE ACCESS	XUCS	FULL	504	13
Filter Predicates				
U.CURSO=233				
TABLE ACCESS	XOCORRENCIAS	FULL	1055	593
Filter Predicates				
O.ANO_LETIVO='2004/2005'				
TABLE ACCESS	XTIPOSALA	FULL	1671	36
Filter Predicates				
S.ANO_LETIVO='2004/2005'				
TABLE ACCESS	XDSD	FULL	27385	27

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			5	77
HASH		GROUP BY	5	77
HASH JOIN			816	76
Access Predicates				
D.ID=ITEM_1				
HASH JOIN			504	49
Access Predicates				
ITEM_2=U.CODIGO				
TABLE ACCESS	YUCS	FULL	504	13
Filter Predicates				
U.CURSO=233				
VIEW	SYS.VW_GBF_10		1671	36
TABLE ACCESS	YTIPOSAULA	FULL	1671	36
Filter Predicates				
S.ANO_LETIVO='2004/2005'				
TABLE ACCESS	YDSD	FULL	27385	27
Other XML				

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			5	74
HASH		GROUP BY	5	74
HASH JOIN			816	73
Access Predicates				
D.ID=ITEM_1				
HASH JOIN			504	46
Access Predicates				
ITEM_2=U.CODIGO				
TABLE ACCESS	ZUCS	FULL	504	13
Filter Predicates				
U.CURSO=233				
VIEW	SYS.VW_GBF_10		1671	33
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED	1671	33
INDEX	TIPOSAULA_IDX_ANOLETIVO	RANGE SCAN	1671	5
Access Predicates				
S.ANO_LETIVO='2004/2005'				
TABLE ACCESS	ZDSD	FULL	27385	27
Other XML				

The usage of the foreign keys codigo, ano_letivo and period improved the cost of the JOIN between yocorrencias and ytiposaula.

The usage of the index tiposaula_idx_anoletivo improved the cost of filtering the ano_letivo in environment z.

4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment X	41
Environment Y	30
Environment Z	27

Query 3 - Which courses (show the code) did have occurrences planned but did not get service assigned in year 2003/2004?

a. Use not in.

1) the SQL query;

```
SELECT o.codigo, o.periodo
FROM xocorrencias o
WHERE o.ano_letivo='2003/2004'
AND (o.codigo, o.periodo) NOT IN
(SELECT t.codigo, t.periodo FROM xtiposaula t WHERE
t.ano_letivo='2003/2004');
```

2) the answer;

CODIGO	PERIODO	CODIGO	PERIODO	CODIGO	PERIODO
1 MEAM1310	3T	32 EM232	2S	48 EIC5126	1S
2 MEST210	2S	33 MEAM1314	3T	49 MEA112	1T
3 MPFCA202	2S	34 CI027	2S	50 EEC2207	2S
4 EIC5124	1S	35 MEA319	3T	51 EIC4223	2S
5 EIC5129	1S	36 MPFCA103	1S	52 EIC4220	2S
6 EEC5022	2S	37 MEAM1312	3T	53 MDI1205	1S
7 EIC4221	2S	38 MPFCA205	2S	54 MPFCA105	1S
8 MPFCA102	1S	39 CI023	1S	55 MPFCA106	1S
9 MPFCA204	2S	40 EIC5125	1S	56 EIC5122	1S
10 EIC4224	2S	41 CI025	2S	57 EIC5127	1S
11 MEM163	2S	42 GEI512	2S	58 MEA414	4T
12 EQ418	2S	43 MEMT110	2S	59 EIC4225	2S
13 MPFCA101	1S	44 MDI1206	1S	60 MPFCA107	1S
14 CI037	2S	45 MEA216	2T	61 EQ411	1S
15 EQ407	2S	46 MEA217	2T	62 MPFCA104	1S
16 MEM181	2S	47 MDI1107	1S	63 MMCCE1220	2T
17 MEM182	2S	48 EIC5126	1S	64 EIC5123	1S
18 MEMT1000	1S	49 MEA112	1T	65 MEMT135	1S
19 MPFCA200	2S	50 EEC2207	2S	66 MDI1108	1S
20 MPFCA100	1S	51 EIC4223	2S	67 EC5200	2S
21 MDI1106	1S	52 EIC4220	2S	68 EQ308	2S
22 MEA415	4T	53 MDI1205	1S	69 MEMT131	2S
23 MEM191	1S	54 MPFCA105	1S	70 MEMT100	1S
24 MPFCA201	2S	55 MPFCA106	1S	71 CI038	2S
25 MDI1105	1S	56 EIC5122	1S	72 MDI1103	1S
26 EIC4222	2S	57 EIC5127	1S	73 MEA215	2T
27 MPFCA206	2S	58 MEA414	4T	74 MEA219	2T
28 MEB205	2S	59 EIC4225	2S	75 MEA320	3T
29 EIC3209	1S	60 MPFCA107	1S	76 MEMT120	2S
30 MEM175	2S	61 EQ411	1S	77 MPFCA203	2S
31 MEMT2000	2S	62 MPFCA104	1S	78 MEB105	2S

3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort;

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			10	629
HASH JOIN		ANTI	10	629
Access Predicates				
AND				
O.CODIGO=T.CODIGO				
O.PERIODO=T.PERIODO				
TABLE ACCESS	XOCORRENCIAS	FULL	1028	593
Filter Predicates				
O.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XTIPOSAULA	FULL	1588	36
Filter Predicates				
T.ANO_LETIVO='2003/2004'				
Other XML				

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			10	63
HASH JOIN		ANTI	10	63
Access Predicates				
AND				
O.CODIGO=T.CODIGO				
O.PERIODO=T.PERIODO				
INDEX	YOCORRENCIAS_PK	FAST FULL SCAN	1028	27
Filter Predicates				
O.ANO_LETIVO='2003/2004'				
TABLE ACCESS	YTIPOSAULA	FULL	1588	36
Filter Predicates				
T.ANO_LETIVO='2003/2004'				
Other XML				

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			10	58
HASH JOIN		ANTI	10	58
Access Predicates				
AND				
O.CODIGO=T.CODIGO				
O.PERIODO=T.PERIODO				
NESTED LOOPS		ANTI	10	58
STATISTICS COLLECTOR				
INDEX	ZOCORRENCIAS_PK	FAST FULL SCAN	1028	27
Filter Predicates				
O.ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED	1588	31
Filter Predicates				
AND				
O.CODIGO=T.CODIGO				
O.PERIODO=T.PERIODO				
INDEX	TIPOSAULA_IDX_ANOLETIVO	RANGE SCAN	1588	5
Access Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED	1588	31
INDEX	TIPOSAULA_IDX_ANOLETIVO	RANGE SCAN	1588	5
Access Predicates				
T.ANO_LETIVO='2003/2004'				
Other XML				

The usage of the primary keys codigo and periodo improved the cost.
The usage of the index tiposaula_idx_anoletivo improved the cost of filtering the ano_letivo in environment z.

4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment X	516
Environment Y	172
Environment Z	115

b. Use external join and is null

1) the SQL query;

```
SELECT o.codigo
FROM xocorrencias o
FULL OUTER JOIN xtiposaula t ON o.ano_letivo=t.ano_letivo AND o.periodo=
t.periodo AND o.codigo=t.codigo
WHERE o.ano_letivo='2003/2004'
AND t.id IS NULL;
```

2) the answer;

CODIGO	CODIGO	CODIGO
1 MDI1105	32 MEMT100	48 MDI1206
2 MDI1106	33 EC5200	49 EIC5123
3 MEA112	34 EIC4220	50 MPFCA204
4 EEC5022	35 MEMT2000	51 EQ418
5 MEA217	36 MPFCA100	52 MEA216
6 MEB205	37 MEA414	53 GEI512
7 CI027	38 MDI1107	54 EIC4224
8 MEAM1310	39 EIC4223	55 MEMT120
9 MDI1108	40 MEAM1314	56 MEM163
10 EQ308	41 EIC5124	57 MPFCA203
11 EM232	42 EEC2207	58 MEMT110
12 MPFCA205	43 MEM182	59 MEMT131
13 EIC5129	44 CI023	60 CI038
14 MMCCE1220	45 MEM175	61 CI025
15 MEMT1000	46 CI037	62 EQ407
16 MEA320	47 MDI1205	63 MEA215
17 MPFCA106	48 MDI1206	64 MEA319
18 MPFCA200	49 EIC5123	65 MEB105
19 EIC4222	50 MPFCA204	66 MEST210
20 EIC4221	51 EQ418	67 MEM181
21 EIC4225	52 MEA216	68 MPFCA202
22 MPFCA104	53 GEI512	69 MEAM1312
23 MEMT135	54 EIC4224	70 MPFCA206
24 EQ411	55 MEMT120	71 MPFCA102
25 MPFCA201	56 MEM163	72 MEA415
26 EIC5122	57 MPFCA203	73 EIC5126
27 EIC5127	58 MEMT110	74 MDI1103
28 MEA219	59 MEMT131	75 MEM191
29 MPFCA107	60 CI038	76 EIC5125
30 EIC3209	61 CI025	77 MPFCA103
31 MPFCA101	62 EQ407	78 MPFCA105

3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort;

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				1 629
FILTER				
Filter Predicates				
T.ID IS NULL				
HASH JOIN		OUTER		1 629
Access Predicates				
AND				
O.ANO_LETIVO=T.ANO_LETIVO(+)				
O.PERIODO=T.PERIODO(+)				
O.CODIGO=T.CODIGO(+)				
TABLE ACCESS	XOCORRENCIAS	FULL		1028 593
Filter Predicates				
O.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XTIPOSAULA	FULL		1588 36
Filter Predicates				
T.ANO_LETIVO(+)='2003/2004'				
Other XML				
SELECT STATEMENT				1028 63
FILTER				
Filter Predicates				
T.ID IS NULL				
HASH JOIN		OUTER		1028 63
Access Predicates				
AND				
O.ANO_LETIVO=T.ANO_LETIVO(+)				
O.PERIODO=T.PERIODO(+)				
O.CODIGO=T.CODIGO(+)				
INDEX	YOCORRENCIAS_PK	FAST FULL SCAN		1028 27
Filter Predicates				
O.ANO_LETIVO='2003/2004'				
TABLE ACCESS	YTIPOSAULA	FULL		1588 36
Filter Predicates				
T.ANO_LETIVO(+)='2003/2004'				
Other XML				
SELECT STATEMENT				1028 58
FILTER				
Filter Predicates				
T.ID IS NULL				
HASH JOIN		OUTER		1028 58
Access Predicates				
AND				
O.ANO_LETIVO=T.ANO_LETIVO(+)				
O.PERIODO=T.PERIODO(+)				
O.CODIGO=T.CODIGO(+)				
INDEX	ZOCORRENCIAS_PK	FAST FULL SCAN		1028 27
Filter Predicates				
O.ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED		1588 31
INDEX	TIPOSAULA_IDX_ANOLETIVO	RANGE SCAN		1588 5
Access Predicates				
T.ANO_LETIVO(+)='2003/2004'				
Other XML				

The usage of the primary keys codigo, periodo and ano_letivo improved the cost in environment y.

The usage of the index tiposaula_idx_anoletivo improved the cost of filtering the ano_letivo in environment z.

4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment X	91
Environment Y	89
Environment Z	84

Query 4 - Who is the professor with more class hours for each type of class, in the academic year 2003/2004? Show the number and name of the professor, the type of class and the total of class hours times the factor.

1) the SQL query;

```
CREATE OR REPLACE VIEW somax AS
SELECT d.nome, d.nr, ds.fator, t.tipo, SUM(ds.horas) horasT
FROM xdsd ds
JOIN xdocentes d ON d.nr=ds.nr
JOIN xtiposaula t ON t.id=ds.id
WHERE t.ano_letivo='2003/2004'
GROUP BY t.tipo, d.nr,d.nome,ds.fator;
```

```
SELECT nr número, nome, tipo, horasT*fator horasxfator, horasT horas FROM
somax WHERE (tipo, horasT) IN (
SELECT tipo, MAX(horasT) maximo FROM somax GROUP BY tipo);
```

2) the answer;

	NÚMERO	NOME	TIPO	HORASXFATOR	HORAS
1	210006	João Carlos Pascoal de Faria	OT	3,5	3,5
2	249564	Cecília do Carmo Ferreira da Silva	TP	26	26
3	208187	António Almerindo Pinheiro Vieira	P	30	30
4	207638	Fernando Francisco Machado Veloso Gomes	T	30,67	30,67

3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort;

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				26 139
HASH JOIN		RIGHT SEMI		26 139
Access Predicates				
AND				
TIPO=TIPO				
HORAST=MAXIMO				
VIEW	SYS.VW_NSQ_1			5 69
HASH		GROUP BY		5 69
VIEW	SOMAX			2570 69
HASH		GROUP BY		2570 69
HASH JOIN				2570 68
Access Predicates				
D.NR=DS.NR				
TABLE ACCESS	XDOCENTES	FULL		939 5
HASH JOIN				2570 63
Access Predicates				
T.ID=DS.ID				
TABLE ACCESS	XTIPOSAULA	FULL		1588 36
Filter Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XDSD	FULL		27385 27
VIEW	SOMAX			2570 69
HASH		GROUP BY		2570 69
HASH JOIN				2570 68
Access Predicates				
D.NR=DS.NR				
TABLE ACCESS	XDOCENTES	FULL		939 5
HASH JOIN				2570 63
Access Predicates				
T.ID=DS.ID				
TABLE ACCESS	XTIPOSAULA	FULL		1588 36
Filter Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XDSD	FULL		27385 27

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			26	139
HASH JOIN		RIGHT SEMI	26	139
Access Predicates				
AND				
TIPO=TIPO				
HORAST=MAXIMO				
VIEW	SYS.VW_NSQ_1		5	69
HASH		GROUP BY	5	69
VIEW	SOMAY		2570	69
HASH		GROUP BY	2570	69
HASH JOIN			2570	68
Access Predicates				
D.NR=DS.NR				
TABLE ACCESS	YDOCENTES	FULL	939	5
HASH JOIN			2570	63
Access Predicates				
T.ID=DS.ID				
TABLE ACCESS	YTIPOSAULA	FULL	1588	36
Filter Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	YDSD	FULL	27385	27
VIEW	SOMAY		2570	69
HASH		GROUP BY	2570	69
HASH JOIN			2570	68
Access Predicates				
D.NR=DS.NR				
TABLE ACCESS	YDOCENTES	FULL	939	5
HASH JOIN			2570	63
Access Predicates				
T.ID=DS.ID				
TABLE ACCESS	YTIPOSAULA	FULL	1588	36
Filter Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	YDSD	FULL	27385	27
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			26	129
HASH JOIN		RIGHT SEMI	26	129
Access Predicates				
AND				
TIPO=TIPO				
HORAST=MAXIMO				
VIEW	SYS.VW_NSQ_1		5	64
HASH		GROUP BY	5	64
VIEW	SOMAZ		2570	64
HASH		GROUP BY	2570	64
HASH JOIN			2570	63
Access Predicates				
D.NR=DS.NR				
TABLE ACCESS	ZDOCENTES	FULL	939	5
HASH JOIN			2570	58
Access Predicates				
T.ID=DS.ID				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED	1588	31
INDEX	TIPOSAULA_IDX_ANOLETIVO	RANGE SCAN	1588	5
Access Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZDSD	FULL	27385	27
VIEW	SOMAZ		2570	64
HASH		GROUP BY	2570	64
HASH JOIN			2570	63
Access Predicates				
D.NR=DS.NR				
TABLE ACCESS	ZDOCENTES	FULL	939	5
HASH JOIN			2570	58
Access Predicates				
T.ID=DS.ID				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED	1588	31
INDEX	TIPOSAULA_IDX_ANOLETIVO	RANGE SCAN	1588	5
Access Predicates				
T.ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZDSD	FULL	27385	27

There is no improvement from environment x to y. This is due to usage of a GROUP BY function with 4 different columns, which makes it very difficult to use correct indexes that could improve the performance.

Between environment y to z, the index tiposaula_idx_anoletivo reduces the cost when creating the VIEW because we are filtering from ano_letivo column.

4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment X	100
Environment Y	98
Environment Z	89

Query 5 - Compare the execution plans (just the environment Z) and the index sizes for the query giving the course code, the academic year, the period, and number of hours of the type 'OT' in the academic years of 2002/2003 and 2003/2004.

a. With a B-tree index on the type and academic year columns of the ZTIPOSAULA table;

b. With a bitmap index on the type and academic year columns of the ZTIPOSAULA table.

1) the SQL query;

a.

```
CREATE INDEX tiposaula_idx ON ztiposaula (tipo, ano_letivo);
```

```
SELECT u.sigla_uc, o.ano_letivo, o.periodo, s.horas_turno
FROM ztiposaula s JOIN zocorrencias o ON o.codigo=s.codigo AND
o.ano_letivo=s.ano_letivo AND o.periodo=s.periodo
JOIN zucs u ON u.codigo = o.codigo
WHERE s.tipo='OT' AND s.ano_letivo LIKE '%2003%';
```

b.

```
DROP INDEX tiposaula_idx;
CREATE BITMAP INDEX tiposaula_idx ON ztiposaula (tipo, ano_letivo);
```

```
SELECT u.sigla_uc, o.ano_letivo, o.periodo, s.horas_turno
FROM ztiposaula s JOIN zocorrencias o ON o.codigo=s.codigo AND
o.ano_letivo=s.ano_letivo AND o.periodo=s.periodo
JOIN zucs u ON u.codigo = o.codigo
WHERE s.tipo='OT' AND s.ano_letivo LIKE '%2003%';
```

2) the answer;

	SIGLA_UC	ANO_LETIVO	PERIODO	HORAS_TURNO
1	EC	2002/2003	2S	0,5
2	EC	2003/2004	2S	0,5

- 3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort;
a.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			10	13
HASH JOIN			10	13
Access Predicates				
U.CODIGO=S.CODIGO				
NESTED LOOPS			10	13
NESTED LOOPS			10	13
STATISTICS COLLECTOR				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED	10	3
INDEX	TIPOSAULA_IDX	RANGE SCAN	10	2
Access Predicates				
S.TIPO='OT'				
Filter Predicates				
S.ANO_LETIVO LIKE '%2003%'				
INDEX	ZUCS_PK	UNIQUE SCAN	1	0
Access Predicates				
U.CODIGO=S.CODIGO				
TABLE ACCESS	ZUCS	BY INDEX ROWID	1	1
TABLE ACCESS	ZUCS	FULL	1	1
Other XML				

b.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			10	46
HASH JOIN			10	46
Access Predicates				
U.CODIGO=S.CODIGO				
NESTED LOOPS			10	46
NESTED LOOPS			10	46
STATISTICS COLLECTOR				
TABLE ACCESS	ZTIPOSAULA	FULL	10	36
Filter Predicates				
AND				
S.TIPO='OT'				
S.ANO_LETIVO LIKE '%2003%'				
INDEX	ZUCS_PK	UNIQUE SCAN	1	0
Access Predicates				
U.CODIGO=S.CODIGO				
TABLE ACCESS	ZUCS	BY INDEX ROWID	1	1
TABLE ACCESS	ZUCS	FULL	1	1
Other XML				

Bitmap indexes work best when the indexed columns have high cardinality, i.e., when the number of distinct values is large compared to the total number of rows in the table. That is not the case in this situation, since the cardinality is 10. For that reason, the B-tree index has a lower cost when comparing with the Bitmap index.

- 4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment Z_A	48
Environment Z_B	140

Query 6 - Select the programs (curso) that have classes with all the existing types.

1) the SQL query;

```
SELECT u.curso
FROM xucs u JOIN xocorrencias o ON u.codigo = o.codigo
JOIN xtuposaula s ON o.codigo=s.codigo AND o.ano_letivo=s.ano_letivo AND
o.periodo=s.periodo
GROUP BY u.curso HAVING COUNT(DISTINCT s.tipo)>4;
```

2) the answer;

	CURSO
1	9461
2	4495
3	9508
4	2021

3) the three execution plans in the three environments, their analysis and of the corresponding estimated effort;

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			6	643
HASH		GROUP BY	6	643
Filter Predicates COUNT(\$vm_col_1)>4				
VIEW	SYS.VM_NWWW_1		404	643
HASH		GROUP BY	404	643
HASH JOIN			21019	642
Access Predicates U.CODIGO=O.CODIGO				
TABLE ACCESS	XUCS	FULL	5396	13
HASH JOIN			21019	629
Access Predicates AND O.CODIGO=S.CODIGO O.ANO_LETIVO=S.ANO_LETIVO O.PERIODO=S.PERIODO				
TABLE ACCESS	XOCORRENCIAS	FULL	21747	593
TABLE ACCESS	XTIPOSADULA	FULL	21019	36
Other XML				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			6	51
HASH		GROUP BY	6	51
Filter Predicates COUNT(\$vm_col_1)>4				
VIEW	SYS.VM_NWWW_1		404	51
HASH		GROUP BY	404	51
HASH JOIN			21019	49
Access Predicates U.CODIGO=S.CODIGO				
TABLE ACCESS	YUCS	FULL	5396	13
TABLE ACCESS	YTIPOSADULA	FULL	21019	36
Other XML				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			6	51
HASH		GROUP BY	6	51
Filter Predicates COUNT(\$vm_col_1)>4				
VIEW	SYS.VM_NWWW_1		404	51
HASH		GROUP BY	404	51
HASH JOIN			21019	49
Access Predicates U.CODIGO=S.CODIGO				
TABLE ACCESS	ZUCS	FULL	5396	13
TABLE ACCESS	ZTIPOSADULA	FULL	21019	36
Other XML				

The usage of the foreign keys codigo, ano_letivo and periodo improved the cost of the JOIN between yocorrencias and ytiposaula.

There is no difference in cost between environment y and z since there is not an index that could improve the performance.

4) if the execution time is measurable, a comparison of the execution times in the three environments.

Environment	Execution Time (ms)
Environment X	91
Environment Y	62
Environment Z	62